

```
[10]: import pandas as pd
import numpy as np
```

```
[11]: import pandas as pd

encodings = ['utf-8', 'latin1', 'ISO-8859-1', 'cp1252']
file_path = "C:\\Users\\MY PC\\Desktop\\Test Jupyter\\machine_learning\\spam.csv"
for encoding in encodings:
    try:
        df = pd.read_csv(file_path, encoding=encoding)
        print(f"File successfully read with encoding: {encoding}")
        break
    except UnicodeDecodeError:
        print(f"Failed to read with encoding: {encoding}")
        continue
if 'df' in locals():
    print("CSV file has been succesfully loaded.")
else:
    print("All encoding attempts failed, unable to read CSV file.")
```

```
Failed to read with encoding: utf-8
File successfully read with encoding: latin1
CSV file has been succesfully loaded.
```

```
[12]: df.sample(5)
```

```
[12]:
```

| | v1 | v2 | Unnamed: 2 | Unnamed: 3 | Unnamed: 4 |
|------|-----|---|------------|------------|------------|
| 2752 | ham | Sat right? Okay thanks... | NaN | NaN | NaN |
| 3338 | ham | Babe !!!! I LOVE YOU !!!! *covers your face in... | NaN | NaN | NaN |
| 4811 | ham | fyi I'm at usf now, swing by the room whenever | NaN | NaN | NaN |
| 4729 | ham | I dont know ask to my brother. Nothing problem... | NaN | NaN | NaN |
| 386 | ham | Customer place i will call you. | NaN | NaN | NaN |

```
[13]: df.shape
```

```
[13]: (5572, 5)
```

```
[14]: # 1. Data Cleaning
# 2. EDA
# 3. Text Processing
# 4. Model Building
# 5. Evaluation
# 6. Improvement
# 7. Website
# 8. Deploy
```

1. Data Cleaning

```
[15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0    v1          5572 non-null   object
1    v2          5572 non-null   object
2    Unnamed: 2   50 non-null     object
3    Unnamed: 3   12 non-null     object
4    Unnamed: 4    6 non-null     object
dtypes: object(5)
memory usage: 217.8+ KB
```

```
[16]: df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], inplace=True)
df.head(5)
```

```
[16]:
```

| | v1 | v2 |
|---|------|---|
| 0 | ham | Go until jurong point, crazy.. Available only ... |
| 1 | ham | Ok lar... Joking wif u oni... |
| 2 | spam | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | ham | U dun say so early hor... U c already then say... |
| 4 | ham | Nah I don't think he goes to usf, he lives aro... |

```
[17]: df.rename(columns={'v1': 'target', 'v2': 'text'}, inplace=True)
df.head(5)
```

```
[17]:
```

| | target | text |
|---|--------|---|
| 0 | ham | Go until jurong point, crazy.. Available only ... |
| 1 | ham | Ok lar... Joking wif u oni... |
| 2 | spam | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | ham | U dun say so early hor... U c already then say... |
| 4 | ham | Nah I don't think he goes to usf, he lives aro... |

```
[23]: from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
df['target'] = encoder.fit_transform(df['target'])
df.head()
```

```
[23]:
```

| | target | text |
|---|--------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... |
| 1 | 0 | Ok lar... Joking wif u oni... |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | 0 | U dun say so early hor... U c already then say... |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... |

```
[19]: # missing values
df.isnull().sum()
```

```
[19]: target    0
text      0
dtype: int64
```

```
[20]: # checking for duplicate values
df.duplicated().sum()
```

```
[20]: 403
```

```
[21]: # remove duplicates
df = df.drop_duplicates(keep='first')
df.duplicated().sum()
```

```
[21]: 0
```

```
[22]: df.shape
```

```
[22]: (5169, 2)
```

2. EDA

```
[24]: df.head()
```


```
[24]:
```

| | target | text |
|---|--------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... |
| 1 | 0 | Ok lar... Joking wif u oni... |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | 0 | U dun say so early hor... U c already then say... |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... |

```
[25]: df['target'].value_counts()
```

```
[25]: target
0      4516
1        653
Name: count, dtype: int64
```

```
[26]: import matplotlib.pyplot as plt
plt.pie(df['target'].value_counts(), labels=['ham', 'spam'], autopct="%0.2f")
plt.show()
```

```
[29]: import nltk
      nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to C:\Users\MY
[nltk_data]   PC\AppData\Roaming\nltk_data...
[nltk_data]   Package punkt is already up-to-date!
```

```
[29]: True
```

```
[30]: df['num_characters'] = df['text'].apply(len) #no of characters
      df.head()
```

```
[30]:
```

| | target | text | num_characters |
|---|--------|---|----------------|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 |

```
[31]: # no of words
df['num_words'] = df['text'].apply(lambda x:len(nltk.word_tokenize(x))) #words count
df.head()
```

```
[31]:
```

| | target | text | num_characters | num_words |
|---|--------|---|----------------|-----------|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 | 13 |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 | 15 |

```
[32]: # sentences
df['num_sentences'] = df['text'].apply(lambda x:len(nltk.sent_tokenize(x))) #sentence count
df.head()
```

```
[32]:
```

| | target | text | num_characters | num_words | num_sentences |
|---|--------|---|----------------|-----------|---------------|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 | 2 |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 | 2 |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 | 2 |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 | 13 | 1 |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 | 15 | 1 |

```
[33]: df[['num_characters', 'num_words', 'num_sentences']].describe()
```

```
[33]:
```

| | num_characters | num_words | num_sentences |
|-------|----------------|-------------|---------------|
| count | 5169.000000 | 5169.000000 | 5169.000000 |
| mean | 78.977945 | 18.455794 | 1.965564 |
| std | 58.236293 | 13.324758 | 1.448541 |
| min | 2.000000 | 1.000000 | 1.000000 |
| 25% | 36.000000 | 9.000000 | 1.000000 |

| | | | |
|------------|------------|------------|-----------|
| 50% | 60.000000 | 15.000000 | 1.000000 |
| 75% | 117.000000 | 26.000000 | 2.000000 |
| max | 910.000000 | 220.000000 | 38.000000 |

```
[34]: # targeting ham
df[df['target']==0][['num_characters', 'num_words', 'num_sentences']].describe()
```

[34]:

| | num_characters | num_words | num_sentences |
|--------------|----------------|-------------|---------------|
| count | 4516.000000 | 4516.000000 | 4516.000000 |
| mean | 70.459256 | 17.123782 | 1.820195 |
| std | 56.358207 | 13.493970 | 1.383657 |
| min | 2.000000 | 1.000000 | 1.000000 |
| 25% | 34.000000 | 8.000000 | 1.000000 |
| 50% | 52.000000 | 13.000000 | 1.000000 |
| 75% | 90.000000 | 22.000000 | 2.000000 |
| max | 910.000000 | 220.000000 | 38.000000 |


```
[35]: # targeting spam
df[df['target']==1][['num_characters', 'num_words', 'num_sentences']].describe()
```

[35]:

| | num_characters | num_words | num_sentences |
|--------------|----------------|------------|---------------|
| count | 653.000000 | 653.000000 | 653.000000 |
| mean | 137.891271 | 27.667688 | 2.970904 |
| std | 30.137753 | 7.008418 | 1.488425 |
| min | 13.000000 | 2.000000 | 1.000000 |
| 25% | 132.000000 | 25.000000 | 2.000000 |
| 50% | 149.000000 | 29.000000 | 3.000000 |
| 75% | 157.000000 | 32.000000 | 4.000000 |
| max | 224.000000 | 46.000000 | 9.000000 |

```
[36]: import seaborn as sns
plt.figure(figsize=(12,6))
sns.histplot(df[df['target']==0]['num_characters'])
sns.histplot(df[df['target']==1]['num_characters'], color='red')
```

```
[36]: <AxesSubplot: xlabel='num_characters', ylabel='Count'>
```




```
[37]: plt.figure(figsize=(12,6))
sns.histplot(df[df['target']==0]['num_words'])
sns.histplot(df[df['target']==1]['num_words'], color='red')
```

```
[37]: <AxesSubplot: xlabel='num words', ylabel='Count'>
```



```
[37]: plt.figure(figsize=(12,6))
sns.histplot(df[df['target']==0]['num_words'])
sns.histplot(df[df['target']==1]['num_words'], color='red')
```


```
[37]: <AxesSubplot: xlabel='num_words', ylabel='Count'>
```



```
[38]: sns.pairplot(df, hue='target')
```

```
[38]: <seaborn.axisgrid.PairGrid at 0x16bb004dae0>
```

[38]: <seaborn.axisgrid.PairGrid at 0x16bb004dae0>

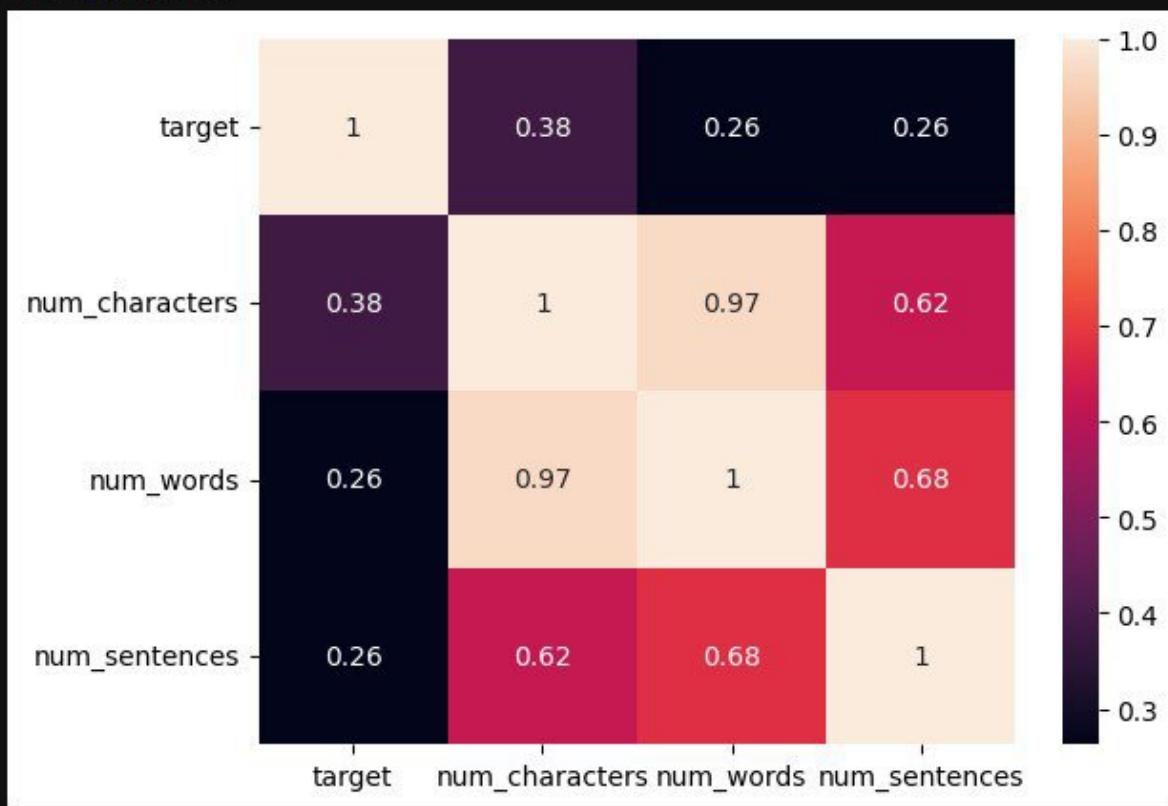


```
[39]: df.dtypes
```

```
[39]: target      int64  
text          object  
num_characters  int64  
num_words      int64  
num_sentences  int64  
dtype: object
```

```
[40]: sns.heatmap(df.corr(numeric_only=True), annot=True)
```

```
[40]: <AxesSubplot: >
```



3. Data Preprocessing

```
[41]: import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
import string

nltk.download('stopwords')

ps = PorterStemmer()

def transform_text(text):
    text = text.lower()
    text = nltk.word_tokenize(text)
    y = []
    for i in text:
        if i.isalnum():
            y.append(i)
    text = y[:]
    y.clear()
    for i in text:
        y.append(ps.stem(i))
    return " ".join(y)

transformed_text = transform_text("I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, k? I've cried enough today")
print(transformed_text)
```

```
[nltk_data] Downloading package stopwords to C:\Users\MY
[nltk_data]   PC\AppData\Roaming\nltk_data...
i gon na be home soon and i do want to talk about thi stuff anymor tonight k i cri enough today
[nltk_data]   Package stopwords is already up-to-date!
```

```
[43]: df['text'][10]
```

```
[43]: "I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, k? I've cried enough today."
```

```
[44]: from nltk.stem.porter import PorterStemmer
ps = PorterStemmer()
ps.stem('walking')
```

```
[44]: 'walk'
```



```
[45]: df['transformed_text'] = df['text'].apply(transform_text)
df.head()
```

```
[45]:
```

| | target | text | num_characters | num_words | num_sentences | transformed_text |
|---|--------|---|----------------|-----------|---------------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 | 2 | go until jurong point crazi avail onli in bugi... |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 | 2 | ok lar joke wif u oni |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 | 2 | free entri in 2 a wkli comp to win fa cup fina... |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 | 13 | 1 | u dun say so earli hor u c already then say |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 | 15 | 1 | nah i do think he goe to usf he live around he... |

```
[46]: from wordcloud import WordCloud
wc = WordCloud(width=500,height=500,min_font_size=10,background_color='white')
spam_wc = wc.generate(df[df['target']==1]['transformed_text'].str.cat(sep=" "))
```


```
[47]: plt.figure(figsize=(15,6))
plt.imshow(spam_wc)
```

```
[47]: <matplotlib.image.AxesImage at 0x16bb30fa6e0>
```



```
[48]: ham_wc = wc.generate(df[df['target']==0]['transformed_text'].str.cat(sep=" "))
plt.figure(figsize=(15,6))
plt.imshow(ham_wc)
```

```
[48]: <matplotlib.image.AxesImage at 0x16bb30a65f0>
```



[49]: df.head()



[49]:


| | target | text | num_characters | num_words | num_sentences | transformed_text |
|---|--------|---|----------------|-----------|---------------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 | 2 | go until jurong point crazi avail onli in bugi... |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 | 2 | ok lar joke wif u oni |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 | 2 | free entri in 2 a wkli comp to win fa cup fina... |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 | 13 | 1 | u dun say so earli hor u c already then say |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 | 15 | 1 | nah i do think he goe to usf he live around he... |

```
[51]: spam_corpus = []
      for msg in df[df['target']==1]['transformed_text'].tolist():
          for word in msg.split():
              spam_corpus.append(word)
```

[52]: len(spam_corpus)

[52]: 14475


```
[53]: from collections import Counter
sns.barplot(x=pd.DataFrame(Counter(spam_corpus).most_common(30))[0],y=pd.DataFrame(Counter(spam_corpus).most_common(30))[1])
plt.xticks(rotation='vertical')
plt.show()
```




```
[55]: ham_corpus = []
for msg in df[df['target']==0]['transformed_text'].tolist():
    for word in msg.split():
        ham_corpus.append(word)
```

```
[56]: len(ham_corpus)
```

```
[56]: 62812
```

```
[57]: from collections import Counter
sns.barplot(x=pd.DataFrame(Counter(ham_corpus).most_common(30))[0],y=pd.DataFrame(Counter(ham_corpus).most_common(30))[1])
plt.xticks(rotation='vertical')
plt.show()
```



```
[58]: # text vectorization
df.head()
```

| | target | text | num_characters | num_words | num_sentences | transformed_text |
|---|--------|---|----------------|-----------|---------------|---|
| 0 | 0 | Go until jurong point, crazy.. Available only ... | 111 | 24 | 2 | go until jurong point crazi avail onli in bugi... |
| 1 | 0 | Ok lar... Joking wif u oni... | 29 | 8 | 2 | ok lar joke wif u oni |
| 2 | 1 | Free entry in 2 a wkly comp to win FA Cup fina... | 155 | 37 | 2 | free entri in 2 a wkli comp to win fa cup fina... |
| 3 | 0 | U dun say so early hor... U c already then say... | 49 | 13 | 1 | u dun say so earli hor u c already then say |
| 4 | 0 | Nah I don't think he goes to usf, he lives aro... | 61 | 15 | 1 | nah i do think he goe to usf he live around he... |

4. Building Model

```
[62]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer  
cv = CountVectorizer()  
tfidf = TfidfVectorizer(max_features=3000)
```

```
[63]: X = tfidf.fit_transform(df['transformed_text']).toarray()
```

```
[64]: #from sklearn.preprocessing import MinMaxScaler  
#scaler = MinMaxScaler()  
#X = scaler.fit_transform(X)  
# appending num_character col to X  
#X = np.hstack((X, df['num_characters'].value.reshape(-1,1)))  
X.shape
```

```
[64]: (5169, 3000)
```

```
[65]: y = df['target'].values
```

```
[66]: from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

```
[67]: from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB  
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score
```

```
[68]: gnb = GaussianNB()  
mnb = MultinomialNB()  
bnb = BernoulliNB()
```

```
[69]: gnb.fit(X_train,y_train)  
y_pred1 = gnb.predict(X_test)  
print(accuracy_score(y_test,y_pred1))  
print(confusion_matrix(y_test,y_pred1))  
print(precision_score(y_test,y_pred1))
```

```
0.8771760154738878
```

```
[[792 104]
```

```
 [ 23 115]]
```

```
0.5251141552511416
```

```
[70]: mnb.fit(X_train,y_train)
y_pred2 = mnb.predict(X_test)
print(accuracy_score(y_test,y_pred2))
print(confusion_matrix(y_test,y_pred2))
print(precision_score(y_test,y_pred2))
```

```
0.9680851063829787
[[896   0]
 [ 33 105]]
1.0
```

```
[71]: bnb.fit(X_train,y_train)
y_pred3 = bnb.predict(X_test)
print(accuracy_score(y_test,y_pred3))
print(confusion_matrix(y_test,y_pred3))
print(precision_score(y_test,y_pred3))
```

```
0.9806576402321083
[[893   3]
 [ 17 121]]
0.9758064516129032
```

```
[72]: from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
```

```
[73]: svc = SVC(kernel='sigmoid', gamma=1.0)
knc = KNeighborsClassifier()
mnb = MultinomialNB()
dtc = DecisionTreeClassifier()
lrc = LogisticRegression(solver='liblinear', penalty='l1')
rfc = RandomForestClassifier(n_estimators=50, random_state=2)
abc = AdaBoostClassifier(n_estimators=50, random_state=2)
bc = BaggingClassifier(n_estimators=50, random_state=2)
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
gbdt = GradientBoostingClassifier(n_estimators=50, random_state=2)
xgb = XGBClassifier(n_estimators=50, random_state=2)
```



```
[74]: clfs = {  
    'SVC' : svc,  
    'KN' : knc,  
    'NB' : mnb,  
    'DT' : dtc,  
    'LR' : lrc,  
    'RF' : rfc,  
    'AdaBoost' : abc,  
    'BgC' : bc,  
    'ETC' : etc,  
    'GBDT' : gbdtd,  
    'xgb' : xgb  
}
```

```
[75]: def train_classifier(clf,X_train,y_train,X_test,y_test):  
    clf.fit(X_train,y_train)  
    y_pred = clf.predict(X_test)  
    accuracy = accuracy_score(y_test,y_pred)  
    precision = precision_score(y_test,y_pred)  
    return accuracy,precision
```

```
[76]: train_classifier(svc,X_train,y_train,X_test,y_test)
```

```
[76]: (0.9816247582205029, 0.983739837398374)
```

```
[77]: accuracy_scores = []  
precision_scores = []  
for name, clf in clfs.items():  
    current_accuracy, current_precision = train_classifier(clf, X_train,y_train,X_test,y_test)  
    print("For ",name)  
    print("Accuracy - ",current_accuracy)  
    print("Precision - ",current_precision)  
    accuracy_scores.append(current_accuracy)  
    precision_scores.append(current_precision)
```

```
For SVC  
Accuracy - 0.9816247582205029  
Precision - 0.983739837398374  
For KN  
Accuracy - 0.90715667311412  
Precision - 1.0  
For NB  
Accuracy - 0.9680851063829787  
Precision - 1.0
```

```
For SVC
Accuracy - 0.9816247582205029
Precision - 0.983739837398374
For KN
Accuracy - 0.90715667311412
Precision - 1.0
For NB
Accuracy - 0.9680851063829787
Precision - 1.0
For DT
Accuracy - 0.9448742746615088
Precision - 0.7913669064748201
For LR
Accuracy - 0.9622823984526112
Precision - 0.9459459459459459
For RF
Accuracy - 0.9709864603481625
Precision - 1.0
```

C:\Users\MY PC\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```
warnings.warn(
For AdaBoost
Accuracy - 0.9700193423597679
Precision - 0.928
For BgC
Accuracy - 0.965183752417795
Precision - 0.9180327868852459
For ETC
Accuracy - 0.9787234042553191
Precision - 0.9833333333333333
For GBDT
Accuracy - 0.960348162475822
Precision - 0.9532710280373832
For xgb
Accuracy - 0.9806576402321083
Precision - 0.9682539682539683
```

```
[82]: In [2]: import pandas as pd
Performance_df = pd.DataFrame({'Algorithm': clfs.keys(), 'Accuracy': accuracy_scores, 'Precision': precision_scores}).sort_values('Precision', ascending=False)
```

```
[83]: performance_df
```

```
[83]:
```

| | Algorithm | Accuracy | Precision |
|----|-----------|----------|-----------|
| 1 | KN | 0.907157 | 1.000000 |
| 2 | NB | 0.968085 | 1.000000 |
| 5 | RF | 0.970986 | 1.000000 |
| 0 | SVC | 0.981625 | 0.983740 |
| 8 | ETC | 0.978723 | 0.983333 |
| 10 | xgb | 0.980658 | 0.968254 |
| 9 | GBDT | 0.960348 | 0.953271 |
| 4 | LR | 0.962282 | 0.945946 |
| 6 | AdaBoost | 0.970019 | 0.928000 |
| 7 | BgC | 0.965184 | 0.918033 |
| 3 | DT | 0.944874 | 0.791367 |


```
[84]: performance_df1 = pd.melt(performance_df, id_vars = "Algorithm")
performance_df1
```

```
[84]:
```

| | Algorithm | variable | value |
|---|-----------|----------|----------|
| 0 | KN | Accuracy | 0.907157 |
| 1 | NB | Accuracy | 0.968085 |
| 2 | RF | Accuracy | 0.970986 |
| 3 | SVC | Accuracy | 0.981625 |
| 4 | ETC | Accuracy | 0.978723 |
| 5 | xgb | Accuracy | 0.980658 |
| 6 | GBDT | Accuracy | 0.960348 |
| 7 | LR | Accuracy | 0.962282 |
| 8 | AdaBoost | Accuracy | 0.970019 |
| 9 | BgC | Accuracy | 0.965184 |

| | | | |
|----|----------|-----------|----------|
| 18 | LR | Precision | 0.945946 |
| 19 | AdaBoost | Precision | 0.928000 |
| 20 | BgC | Precision | 0.918033 |
| 21 | DT | Precision | 0.791367 |

```
[85]: sns.catplot(x="Algorithm", y="value", hue="variable", data=performance_df1, kind='bar', height=5)
plt.ylim(0.5,1.0)
plt.xticks(rotation='vertical')
plt.show()
```




```
[87]: ax_features parameter of Tfidf
performance_df = pd.DataFrame({'Algorithm': clfs.keys(), 'Accuracy_max_ft_3000': accuracy_scores, 'Precision_num_chars': precision_scores}).sort_values('Precision_num_chars', ascending=False)
```

```
[88]: new_df = performance_df.merge(temp_df, on='Algorithm')
```

```
[89]: new_df_scaled = pd.DataFrame({'Algorithm': clfs.keys(), 'Accuracy_num_chars': accuracy_scores, 'Precision_num_chars': precision_scores}).sort_values('Precision_num_chars', ascending=False)
```

```
[90]: new_df_scaled.merge(temp_df, on='Algorithm')
```

```
[90]:
```


| | Algorithm | Accuracy | Precision | Accuracy_max_ft_3000_x | Precision_num_chars_x | Accuracy_max_ft_3000_y | Precision_num_chars_y | Accuracy_num_chars | Precision_num_chars |
|----|-----------|----------|-----------|------------------------|-----------------------|------------------------|-----------------------|--------------------|---------------------|
| 0 | KN | 0.907157 | 1.000000 | 0.907157 | 1.000000 | 0.907157 | 1.000000 | 0.907157 | 1.000000 |
| 1 | NB | 0.968085 | 1.000000 | 0.968085 | 1.000000 | 0.968085 | 1.000000 | 0.968085 | 1.000000 |
| 2 | RF | 0.970986 | 1.000000 | 0.970986 | 1.000000 | 0.970986 | 1.000000 | 0.970986 | 1.000000 |
| 3 | SVC | 0.981625 | 0.983740 | 0.981625 | 0.983740 | 0.981625 | 0.983740 | 0.981625 | 0.983740 |
| 4 | ETC | 0.978723 | 0.983333 | 0.978723 | 0.983333 | 0.978723 | 0.983333 | 0.978723 | 0.983333 |
| 5 | xgb | 0.980658 | 0.968254 | 0.980658 | 0.968254 | 0.980658 | 0.968254 | 0.980658 | 0.968254 |
| 6 | GBDT | 0.960348 | 0.953271 | 0.960348 | 0.953271 | 0.960348 | 0.953271 | 0.960348 | 0.953271 |
| 7 | LR | 0.962282 | 0.945946 | 0.962282 | 0.945946 | 0.962282 | 0.945946 | 0.962282 | 0.945946 |
| 8 | AdaBoost | 0.970019 | 0.928000 | 0.970019 | 0.928000 | 0.970019 | 0.928000 | 0.970019 | 0.928000 |
| 9 | BgC | 0.965184 | 0.918033 | 0.965184 | 0.918033 | 0.965184 | 0.918033 | 0.965184 | 0.918033 |
| 10 | DT | 0.944874 | 0.791367 | 0.944874 | 0.791367 | 0.944874 | 0.791367 | 0.944874 | 0.791367 |

```
[91]: # voting Classifier
svc = SVC(kernel='sigmoid', gamma=1.0, probability=True)
mnb = MultinomialNB()
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
from sklearn.ensemble import VotingClassifier
```

```
[91]: # voting Classifier
svc = SVC(kernel='sigmoid', gamma=1.0, probability=True)
mnb = MultinomialNB()
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
from sklearn.ensemble import VotingClassifier
```

```
[92]: voting = VotingClassifier(estimators=[('svm', svc), ('nb', mnb), ('et', etc)], voting='soft')
```

```
[93]: voting.fit(X_train, y_train)
```



```
[94]: y_pred = voting.predict(X_test)
print("Accuracy", accuracy_score(y_test, y_pred))
print("Precision", precision_score(y_test, y_pred))
```

```
Accuracy 0.9825918762088974
Precision 0.9918032786885246
```

```
[95]: # Applying Stacking
estimators=[('svm', svc), ('nb', mnb), ('et', etc)]
final_estimator=RandomForestClassifier()
```

```
[96]: from sklearn.ensemble import StackingClassifier
clf = StackingClassifier(estimators=estimators, final_estimator=final_estimator)
```

```
[97]: clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print("Accuracy", accuracy_score(y_test, y_pred))
print("Precision", precision_score(y_test, y_pred))
```

```
Accuracy 0.9816247582205029
Precision 0.9541984732824428
```

```
[98]: import pickle
```

```
[99]: pickle.dump(tfidf, open('vectorizer.pkl','wb'))  
pickle.dump(mnb, open('model.pkl','wb'))
```


```
[107]: import pickle  
from sklearn.feature_extraction.text import TfidfVectorizer  
from sklearn.naive_bayes import MultinomialNB  
  
# Sample text data and corresponding labels (replace with actual data)  
X_train = ["Sample text 1", "Sample text 2", "Sample Text 3"]  
y_train = [0,1,0] # 0 for negative and 1 for positive  
  
# create and train the TF-IDF vectorizer  
tfidf = TfidfVectorizer(lowercase=True, stop_words='english')  
X_train_tfidf = tfidf.fit_transform(X_train)  
  
# create and train the Naive Bayes classifier  
mnb = MultinomialNB()  
mnb.fit(X_train_tfidf, y_train)  
  
# Save the trained TF-IDF vectorizer and Naive Bayes model to files  
with open('vectorizer.pkl','wb') as vectorizer_file:  
    pickle.dump(tfidf, vectorizer_file)  
with open('model.pkl','wb') as model_file:  
    pickle.dump(mnb, model_file)
```

```
[ ]:
```



```
1  import streamlit as st
2  import pickle
3  tfidf = pickle.load(open('vectorizer.pkl', 'rb'))
4  model = pickle.load(open('model.pkl', 'rb'))
5
6  import nltk
7  from nltk.corpus import stopwords
8  from nltk.stem import PorterStemmer
9  import string
10
11  nltk.download('stopwords')
12
13  ps = PorterStemmer()
14
15  def transform_text(text):
16      text = text.lower()
17      text = nltk.word_tokenize(text)
18      y = []
19      for i in text:
20          if i.isalnum():
21              y.append(i)
22      text = y[:]
23      y.clear()
24      for i in text:
25          y.append(ps.stem(i))
26      return " ".join(y)
27
28  st.title("Email Spam Classifier")
29  input_sms = st.text_area("Enter message")
30
31  if st.button('Predict'):
32      transformed_data = transform_text(input_sms)
33      vector_input = tfidf.transform([transformed_data])
34      result = model.predict(vector_input)[0]
35      if result == 1:
36          st.header("Spam")
37      else:
38          st.header("Not Spam")
```


PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL JUPYTER: VARIABLES



streamlit
streamlit
streamlit
streamlit

Install the latest PowerShell for new features and improvements! <https://aka.ms/PSWindows>

PS C:\Users\MY PC\Desktop\Test Jupyter\machine_learning> streamlit run app.py

You can now view your Streamlit app in your browser.

Local URL: <http://localhost:8504>

Network URL: <http://192.168.43.130:8504>





localhost:8503



Deploy



Email Spam Classifier

Enter message

Hy, this is an important [email](#).

Predict

Not Spam